

AI Critique

Across HW03 the AI accelerated scaffolding, but it repeatedly produced work that looked finished yet was wrong until a human checked it against ground truth. The clearest failure was the usability study: the AI first generated seven *simulated* personas and even wrote "Persona and behavior were simulated" into every observation note, which directly contradicts the assignment's demand for seven real participants. It took explicit human redirection to replace this with genuine recorded sessions. The AI also left internal contradictions it had introduced — one participant's header said order #7 "Hủy" while the event log still read "Mã 6, ngày 26/7" — because it edited fields in isolation without re-reading the whole file. Most tellingly, when a leftover mobile-only answer had to be removed, the AI immediately *fabricated* a replacement quote for the participant; that is an integrity failure, corrected only because it was caught. In Task 3 the AI first offered WebKit-on-Windows as "Safari" evidence, which the rubric forbids, so a real Safari-on-iOS run was needed. It also obediently deleted a required masked-contact column simply because it was asked, without flagging that the specification mandates it.

The common cause is that the AI optimises for producing a complete-looking artifact, so it fills gaps with assumptions or invented data and trusts its own prior output. The principle I learned is to treat every AI "done" as a claim to be verified against the real system and the actual requirements, to keep the AI grounded in collected evidence, and never to let it fabricate what only a human or the SUT can legitimately supply.